
Group no C38

Vatsal Gupta 200101105

Sweeya Reddy 200101079

Pranshu Kandoi 200101086

Assignment-2 OS344

PART A

Here following system calls need to be implemented **getNumProc()**, **getMaxPid()**, **getProcInfo(pid, &processInfo)**, **set_burst_time(n)** and **get_burst_time()**. Note defs and processInfo.h files are provided later containing the processInfo data structure. Also in params.h the number of cpu is also set to 1. Also note that to get current process info and pid 2 system calls are also added which will be used later in the assignment.

So for the purpose of implementing these system calls following files are changed:

1. **We will create system calls** and make their corresponding links. Indexes for the system calls that we will implement are defined in **syscall.h**. Thus it contains the mapping from system call name to system call number.

```
1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat   8
10 #define SYS_chdir   9
11 #define SYS_dup    10
12 #define SYS_getpid 11
13 #define SYS_sbrk   12
14 #define SYS_sleep  13
15 #define SYS_uptime 14
16 #define SYS_open   15
17 #define SYS_write  16
18 #define SYS_mknod  17
19 #define SYS_unlink 18
20 #define SYS_link   19
21 #define SYS_mkdir  20
22 #define SYS_close  21
23 #define SYS_getNumProc 22
24 #define SYS_getMaxPID 23
25 #define SYS_getProcInfo 24
26 #define SYS_set_burst_time 25
27 #define SYS_get_burst_time 26
28
```

-
2. **syscall.c** contains helper functions to parse call arguments and pointers to the actual system call implementations.

```
133 [SYS_close] sys_close,  
134 [SYS_getNumProc] sys_getNumProc,  
135 [SYS_getMaxPID] sys_getMaxPID,  
136 [SYS_getProcInfo] sys_getProcInfo,  
137 [SYS_set_burst_time] sys_set_burst_time,  
138 [SYS_get_burst_time] sys_get_burst_time,  
139 };  
140  
.04 extern int sys_write(void);  
.05 extern int sys_uptime(void);  
.06 extern int sys_getNumProc(void);  
.07 extern int sys_getMaxPID(void);  
.08 extern int sys_getProcInfo(void);  
.09 extern int sys_set_burst_time(void);  
.10 extern int sys_get_burst_time(void);  
.11
```

3. **usys.S** contains a list of system calls exported by the kernel.

```
31 SYSCALL(uptime)  
32 SYSCALL(getNumProc)  
33 SYSCALL(getMaxPID)  
34 SYSCALL(getProcInfo)  
35 SYSCALL(set_burst_time)  
36 SYSCALL(get_burst_time)|  
37
```

4. **Proc.h** contains the struct proc structure. In this number of context fields, burst time and running time fields have been added. Also system call has been defined in proc.h.

```

37 // Per-process state
38 struct proc {
39     uint sz;                      // Size of process memory (bytes)
40     pde_t* pgdir;                // Page table
41     char *kstack;                // Bottom of kernel stack for this process
42     enum procstate state;        // Process state
43     int pid;                     // Process ID
44     struct proc *parent;         // Parent process
45     struct trapframe *tf;        // Trap frame for current syscall
46     struct context *context;    // swtch() here to run process
47     void *chan;                  // If non-zero, sleeping on chan
48     int killed;                 // If non-zero, have been killed
49     struct file *ofile[NFILE];   // Open files
50     struct inode *cwd;          // Current directory
51     char name[16];              // Process name (debugging)
52     int nocs;                   // Number of context switches
53     int burst_time;             // Burst time duhhh
54     int rt;                     // running time
55 };
56
57 int getNumProc_system(void);
58 int getMaxPID_system(void);
59 struct processInfo getProcInfo_system(int pid);
60 int set_burst_time_system(int burst_time);
61 int get_burst_time_system();
62

```

5. **user.h** contains system call definitions in xv6.

```

26 int uptime(void);
27 int getNumProc(void);
28 int getMaxPID(void);
29 int getProcInfo(int, struct processInfo*);
30 int set_burst_time(int);
31 int get_burst_time();
32

```

6. In makefile, in **EXTRA** the following test file name is added and also in **UPROGS** for linking our test files and corresponding system calls.

```
5/  
68     UPROGS=\  
69     |_cat\  
70     |_echo\  
71     |_forktest\  
72     |_grep\  
73     |_init\  
74     |_kill\  
75     |_ln\  
76     |_ls\  
77     |_mkdir\  
78     |_rm\  
79     |_sh\  
80     |_stressfs\  
81     |_usertests\  
82     |_wc\  
83     |_zombie\  
84     |_getNumProcTest\  
85     |_getMaxPIDTest\  
86     |_getProcInfoTest\  
87     |_burstTimeTest\  
88     |_ioProcTester\  
89     |_cpuProcTester\  
90     |_test_scheduler\  
91
```

```
ll.c\  
bie.c getNumProcTest.c getMaxPIDTest.c getProcInfoTest.c burstTimeTest.c test_scheduler.c ioProcTester.c cpuProcTester.c
```

7. **sysproc.c** contains the implementations of process related system calls.

```
93
94     int
95     sys_getNumProc(void)
96     {
97         return getNumProc_system();
98     }
99
100    int
101    sys_getMaxPID(void)
102    {
103        return getMaxPID_system();
104    }
105
106    int
107    sys_getProcInfo(void){
108        int pid;
109        struct processInfo *info;
110        argptr(0,(void *)&pid, sizeof(pid));
111        argptr(1,(void *)&info, sizeof(info));
112
113        struct processInfo tempInfo = getProcInfo_system(pid);
114
115        if(tempInfo.ppid == -1) return -1;
116        info->ppid = tempInfo.ppid;
117        info->psize = tempInfo.psize;
118        info->numberContextSwitches = tempInfo.numberContextSwitches;
119        return 0;
120    }
121
```

```
int
sys_set_burst_time(void)
{
    int burst_time;
    argptr(0,(void *)&burst_time, sizeof(burst_time));

    return set_burst_time_system(burst_time);
}

int
sys_get_burst_time(void)
{
    return get_burst_time_system();
}
```

8. **proc.c**, the assist functions used by system calls are defined in the proc.c

```

1 int getNumProc_system(void){
2
3     int num_processes;
4     struct proc *p;
5
6     num_processes = 0;
7     acquire(&ptable.lock);
8     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
9         if(p->state != UNUSED){
10             num_processes++;
11         }
12     }
13     release(&ptable.lock);
14
15     return num_processes ;
16 }
17
18 int getMaxPID_system(void){
19
20     int max_pid;
21     struct proc *p;
22
23     max_pid = 0;
24     acquire(&ptable.lock);
25     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
26         if(p->state != UNUSED){
27             if(p->pid > max_pid){
28                 max_pid = p->pid;
29             }
30         }
31     }
32     release(&ptable.lock);
33
34     return max_pid;
35 }
36
37 int set_burst_time_system(int burst_time)
38 {
39     struct proc *p = myproc();
40     p->burst_time = burst_time;
41     acquire(&quantlock);
42     if(burst_time < quant){
43         quant = burst_time;
44     }
45     release(&quantlock);
46     yield();
47
48     return 0;
49 }
50
51 int get_burst_time_system()
52 {
53     struct proc *p = myproc();
54
55     return p->burst_time;
56 }
57

```

```

1 struct processInfo getProcInfo_system(int pid){
2
3     struct proc *p;
4     struct processInfo temporary_info = {-1,0,0};
5
6     acquire(&ptable.lock);
7     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
8         if(p->state != UNUSED){
9             if(p->pid == pid) {
10                 temporary_info.ppid = p->parent->pid;
11                 temporary_info.psize = p->sz;
12                 temporary_info.numberContextSwitches = p->nocs;
13                 release(&ptable.lock);
14                 return temporary_info;
15             }
16         }
17     }
18     release(&ptable.lock);
19
20     return temporary_info;
21 }
22

```

Part 1

For getting the number of processes and the maximum pid we will simply acquire a lock on the process table and loop through the table and for each process which is not unused we will add to the number of processes and also see if its the maximum pid process. This whole thing is done in the assist function in proc.c in **getNumProc_system** and **getMaxPID_system** and then in sysproc.c the following function is used. These function implementations are shown above.

```
SeaBIOS (version 1.13.0-1ubuntu1.1)

iPXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CA10+1FECCA10 CA00

Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ getNumProcTest
The total number of active processes in the system is 3
$ getMaxPIDTest
The maximum PID of all active processes in the system is 4
$
```

We can note that since 2 system processes are always running thus total process running are 3 as 2 system processes are there and one the current **getNumProcTest** and also since the **getMaxPIDTest** is the process with PID 4 and is the max among all active process as right now this one is also running thus maxPID is 4.

Part 2

For implementing **getProcInfo** we are providing 2 arguments one is the pid of the required process and other is the space where we need to store that info thus similar to previous a lock is maintained on the process table and is searched for the current process if the current process is found then lock is released and the process info is saved to the location and if not found then default value is returned by the helper function in proc.c also system call in sysproc.c uses this to get the required info. Also the process parent is set to default parent and number of context switches and burst time are set to 0 by default in proc.c. Also we set the **default parent pid to -2** which helps us to determine whether the current process has a parent or not. For calculating the number of context switches we will simply **increase the value of the nocs variable when the process is being scheduled by the scheduler**.

```
Booting from Hard Disk...
xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ getProcInfoTest 1
PPID: No Parent Process
Size: 16384
Number of Context Switches: 17
$ getProcInfoTest 2
PPID: 1
Size: 20480
Number of Context Switches: 22
```

We can see from the output that the process with pid 1 has no parent and the process with pid 2 has a parent with pid 1. Also the context switches are shown in the above picture.

Part 3

For this part we simply need to add a burst time field in the proc data structure and created two system call to set this value and get this value as shown above in sysproc.c and proc.c . sysproc.c take what value to be set as burst time and helper function in proc.c set that burst time. Note after burst time is set yield() is called as we need to remove the process from running as its burst time has changed thus it needs to be evaluated by the scheduling algorithm again. Also to get the current process we can simply use myproc() function.

```
c burstTimeTest.c > ...
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int
6 main(void){
7     printf(1,"Burst time is set to 10\n");
8     set_burst_time(10);
9     printf(1, "Getting the value of burst time %d\n", get_burst_time());
10    exit();
11 }
```

```
Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ burstTimeTest
Burst time is set to 10
Getting the value of burst time 10
$ 
```

From the output we clearly see that the current process burst time is set to 10.

PART B

SJF (Shortest Job First) Scheduler

Firstly for implementing the shortest job scheduler without the round robin scheduler in trap.c the following line needs to be commented as to prevent any process to be evicted based on time quantum. Thus commenting this line will prevent forced preemption of the process and scheduling is done based on the shortest job first not round robin.

```

102
103     // Force process to give up CPU on clock tick.
104     // If interrupts were on while locks held, would need to check nlock.
105     if(myproc() && myproc()->state == RUNNING &&
106        tf->trapno == T_IRQ0+IRQ_TIMER){
107         yield();
108     }
109

```

For the purpose of implementing the scheduling algorithm a **min heap** is maintained using an array of all the processes whose state is RUNNABLE and proper locking is handled to add and remove processes from the queue. Thus functions are also written to handle various min heap operations. Also whenever a function is made runnable then we add this to our min heap or priority queue.

<pre> // //new SJF SCHEDULER acquire(&phtable.lock); if((p = extractMin()) == 0){ release(&phtable.lock); continue; } if(p->state!=RUNNABLE) { release(&phtable.lock); continue; } c->proc = p; switchuvm(p); p->state = RUNNING; swtch(&(c->scheduler), p->context); (p->nocs)++; switchkvm(); c->proc = 0; release(&phtable.lock); </pre>	<pre> struct { struct spinlock lock; int siz; struct proc* proc[NPROC+1]; } pqueue; </pre>
--	--

Note here firstly using extract min the minimum burst time process is extracted and then scheduled thus a shortest job first scheduler. Also note that in the pqueue structure spinlock is there, size of the queue is there and an array is used which act as a min heap.

Time complexity and corner cases

Running complexity of the algorithm is $O(\log n)$ as the extract min operation in the priority queue is of $O(\log n)$ and then the remaining operation is of order 1. Note for handling corner cases firstly if the queue is full no new process is inserted into the queue and it returns as it is. If the queue is empty then extract min returns 0 and thus the scheduler does not schedule any processes. Also if the process state is already runnable then it is already in the queue and thus it is not added in the queue. Also if a process is finished that is when its state is ZOMBIE then also it is removed from the queue. Also when using priority queue a locking is done properly and also only those processes whose state are runnable are scheduled otherwise is not scheduled.

Testing

For testing purposes a test_scheduler.c file is created which take an integer n as input and for half of that uses for io bound processes and half of that uses as cpu bound processes. For a cpu bound process a double for loop is there and for a io bound process sleep(1) is written in for loop for 50 times. Thus calling test_scheduler gives following output

```
p: size 1000 nblocks 941 ninode 200 nlog 30 logstart 2 inode
init: starting sh
test_scheduler 14
    PID      Type     Burst Time      Context Switches
    ---      ----     -----      -----
    11       CPU       2              2
    13       CPU       3              2
    7        CPU       3              2
    17       CPU       5              2
    15       CPU       7              2
    9        CPU      10             2
    5        CPU      10             2
    8        I/O       1              102
    6        I/O       1              102
    10      I/O       3              102
    4        I/O       7              102
    12      I/O       9              102
    14      I/O      10             102
    16      I/O      10             102
```

Note here firstly that the processes are scheduled based on the burst time and also on the type of the processes. Note firstly that the CPU bound processes completes first and then the IO bound processes as in IO bound processes because of sleep it takes more time to complete. Also note in CPU and also IO bound processes processes are completed based on the burst time. Also note that there are 2 context switches in case of CPU bound processes as 1st for when the process is scheduled and when its burst time is set. Note in case of IO bound processes because of 100 sleep additional 100 context switches are there.

Given below is that of the default round robin scheduler and we can see that the scheduling happens mostly on the basis PID values of the processes and also context switches in this case are very high as compared to the shortest job first.

```
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inode
init: starting sh
$ test_scheduler 14
      PID      Type    Burst Time    Context Switches
      ---      ----    -----      -----
      5        CPU      10            34
      7        CPU      3             33
      9        CPU      10            33
     11        CPU      2             33
     15        CPU      7             34
     13        CPU      3             35
     17        CPU      5             35
      4        I/O      7             102
      6        I/O      1             103
      8        I/O      1             102
     10        I/O      3             102
     12        I/O      9             102
     14        I/O      10            102
     16        I/O      10            102
$ □
```

For another program we can see the same difference

Round robin

```
init: starting sh
test_scheduler 14
      PID      Type    Burst Time    Context Switches
      ---      ----    -----      -----
      7        CPU      3             73
      5        CPU      10            74
     11        CPU      2             73
      9        CPU      10            75
     15        CPU      7             73
     13        CPU      3             75
     17        CPU      5             75
      4        I/O      7             202
      6        I/O      1             202
      8        I/O      1             202
     10        I/O      3             202
     12        I/O      9             202
     14        I/O      10            202
     16        I/O      10            202
$ □
```

shortest job first

```
size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inode
: starting sh
st_scheduler 14
      PID      Type    Burst Time    Context Switches
      ---      ----    -----      -----
     11        CPU      2             2
     13        CPU      3             2
      7        CPU      3             2
     17        CPU      5             2
     15        CPU      7             2
      9        CPU      10            2
      5        CPU      10            2
      6        I/O      1             202
      8        I/O      1             202
     10        I/O      3             202
      4        I/O      7             202
     12        I/O      9             202
     14        I/O      10            202
     16        I/O      10            202
$ □
```

Hybrid Scheduler

For the purpose of a hybrid scheduler we firstly need to maintain 2 priority queues. Firstly job are scheduled on the basis of the first queue in which all the runnable processes are there and then after a fixed time quantum it is preempted and put to the second queue. After the 1st queue is empty it means that the 1st round is completed and all the remaining processes in the second are again put to the first queue and again scheduled based on the shortest burst time. And thus we are able to achieve a hybrid scheduler. Also note that the time quantum variable is modified

based on the burst time, that is if the burst time is less than the time quantum its value is decreased to that value.

Finally in trap.c interrupt code need to be modified as shown below

```
// Force process to give up CPU on clock tick.  
// If interrupts were on while locks held, would need to ch  
if(myproc() && myproc()->state == RUNNING &&  
    tf->trapno == T_IRQ0+IRQ_TIMER){  
    [myproc()->rt]++;  
    if(myproc()->burst_time != 0){  
        if(myproc()->burst_time == myproc()->rt)  
            exit();  
    }  
    if((myproc()->rt)%quant == 0) new_yield();  
    // yield();  
}
```

Note here rt is the running time variable keeping track of how long the process is executed and if it becomes equal to burst time then that process is completed exit command is executed. If not then new_yield is executed and the current process is put to the other queue.

Thus the scheduler code looks like

```
for(;;){  
    // Enable interrupts on this processor.  
    sti();  
    acquire(&ptable.lock);  
    //Hybrid scheduler  
    if(isEmpty()){  
        if(isEmpty2())  
            goto label;  
    }  
    while(!isEmpty2()){  
        if((p = extractMin()) == 0){release(&ptable.lock);break;}  
        insert(p,0);  
    }  
}  
label:  
    if((p = extractMin()) == 0){  
        release(&ptable.lock);  
        continue;  
    }  
    if(p->state!=RUNNABLE){  
        release(&ptable.lock);  
        continue;  
    }  
  
    c->proc = p;  
    switchvmm(p);  
    p->state = RUNNING;  
    (p->nocs)++;  
  
    swtch(&(c->scheduler), p->context);  
    switchkvm();  
    c->proc = 0;  
    release(&ptable.lock);
```

Note firstly that the main queue is empty then all the remaining processes from the other queue are put into the main queue then similar to the previous shortest job first scheduler, jobs are scheduled. Other functions of the other priority queue are identical to that of the main priority queue. Also a proper locking system is used.

hit: starting sh test_scheduler 14				
PID	Type	Burst Time	Context Switches	
6	I/O	11	52	
8	I/O	36	52	
10	I/O	126	52	
7	CPU	132	10	
13	CPU	135	10	
17	CPU	222	10	
15	CPU	338	10	
9	CPU	493	10	
5	CPU	500	10	
4	I/O	318	52	
12	I/O	411	52	
16	I/O	466	52	
14	I/O	481	52	

From the hybrid scheduler we can clearly see that the number of content switches have increased in case of cpu bound processes as compared to shortest job scheduling as process are now been preempted more due to round robin scheduling but in case of IO bound processes since most of the time they are in sleep thus difference is not well observed and context switches are thus same as that of sjf scheduler also note that process are completed based on the burst time of the cpu and io type processes.